

**Hey look we have a cluster now
isn't that cool what does it do does it do stuff yes it
does it's cool**

Doing cluster stuff for fun and ~~profit~~ science

Lukas Burk

Leibniz Institute for Prevention Research and Epidemiology — BIPS

2026-06-18

BioWiMium

We have a new cluster

You might have *heard*

What we cover today:

- What is this cluster thing
- How does it work (using R / Python, Slurm)
- How do **you** do stuff on it in general

We don't cover

- How to Linux (from the command line)
- How to use batchtools (or targets)
- Everything you ever need to know



I know it's noisy and no I don't know when that is fixed

What is a cluster?

And why should you care

Many ways to shave a yak, and now you have one more!

Your PC/laptop

- You know it
- Shorter tasks
- Not so powerful
- Convenience: 🖱️

Workstation / Server

- Bertha / new GePaRD server
- Run looong tasks
- 1 powerful computer
- Shared
- Convenience: Mixed

HPC Cluster

- Many computers connected
- Access via *head node*
- Shared (scheduling: *Slurm*)
- Convenience: You'll get there

What does it look like?

Kind of like this



The HPC workflow

It's a freight truck, not a bicycle

You don't hop on the cluster for a quick ride, you **submit work** to it

What you are used to

- Open RStudio / VSCode / Positron
- Write code, run it, see results
- Everything is *interactive*
- One computer does it all

How HPC works

- Prepare your code on the *head node*
- **Submit** it as a job (it gets *queued*)
- Go make coffee ☕
- Come back, check results
- Debug if needed, resubmit

💡 There's a learning curve, but the payoff is huge

HPC Workflow

The typical cycle




5

1. Log in (land on **head node**)
2. Load software (`module load R/4.5.3`)
3. Move to your project (`git helps`)
4. Submit your **job(s)** via `sbatch` (or `salloc` for interactive work)
5. Monitor with `squeue --me`
6. Log out, check back later
7. If jobs failed, debug & resubmit
8. Repeat until **insane** done


Cluster topology

Partitions



Cluster topology

It only looks weird because it is



The head node

Your gateway

The head node is where you land when you log in

Do ✓

- Install R / Python packages
- Edit scripts and code
- Submit and monitor jobs
- Use `git`

Don't ✗


- Run heavy computations
- Launch long-running scripts
- Load huge datasets into memory

⚠ The head node is shared by *everyone*. Heavy processes will be *terminated*.

Terminology

Just to make sure you're confused

- The cluster has 12 compute **nodes** (self-contained computer)
 - Each node has 1 **CPU** (physical processor)
 - Each processor is made up of 96 **cores** (independent units)
 - Each core can juggle 2 **threads** (e.g. R processes)
- Terminology here weird / confusing because: *history* 🙄
 - Until the 90s: **1** CPU = **1** core = **1** thread
 - Then: “Hyperthreading”: **1** CPU = **1** core = **2** threads
 - Then: “Dual core”: **1** CPU = **2** cores = **4** threads
 - **Now**: **1** CPU = **96** cores = **192** threads



⚠ **Slurm**: You need 10 **threads** → you request 10 “cpus”

But what is Slurm even

- Slurm is the job management system (see slurm.schedmd.com)
- Every CPU *core*, every Byte of *RAM* on the compute nodes is *kept track of*

→ You can only use resources allocated to you

Things that *do not* happen on a cluster:

- “My job took too much RAM so your jobs got killed alongside mine sorry”
- “I used 1000 threads but only meant to use 10 sorry your jobs are smothered”

→ It’s impossible to “endanger” other people’s compute jobs

Logging in

Different environments for different tasks

11

Connect to the cluster **head node** via SSH in one of two ways:

Terminal

- Definitely cool and normal
- 1337 h4xX0r feeling

```
burk hnode ~/projects
```

Permissions	User	Group	Size	Date Modified	Name
wxr-xr-x	burk	burk	4.0K	2026-04-02 19:35	📁 bips-cluster-demos/
wxr-xr-x	burk	burk	4.0K	2026-03-19 00:12	📁 paper_2023_survival_benchmark/
wxr-xr-x	burk	burk	4.0K	2026-03-05 00:12	📁 paper_2025_xplainfi_benchmark/
wxr-xr-x	burk	burk	4.0K	2026-03-31 13:23	📁 reduction-techniques-benchmark/
wxr-xr-x	burk	burk	4.0K	2026-03-31 15:50	📁 reduction-techniques-benchmark-update/
wxr-xr-x	burk	burk	4.0K	2026-03-17 20:13	📁 resolve-pcc-symptoms-multi-analyst/
wxr-xr-x	burk	burk	4.0K	2026-03-18 01:19	📁 resolve-pcc-xai/
wxr-xr-x	burk	burk	4.0K	2026-04-01 18:31	📁 survival-benchmark/
wxr-xr-x	burk	burk	4.0K	2026-03-20 21:15	📁 xplainfi-benchmark/

Positron (VScode)

- Uses SSH for communication with head node
- GUI on your own device (PC, laptop)
- Terminal / R console runs **on head node**
- All files you see are on the head node
- 😎 Edit files conveniently

The first thing you see when you log in

Trying to point you in the right direction



12

A large ASCII art graphic of a computer monitor, rendered in light blue dashed lines on a dark background. The screen area of the monitor contains the text 'BIPS cluster' in a similar dashed font.

Welcome to the BIPS HPC cluster.

Dashboard: <http://srvcluster.bips.de>
Docs: <https://cluster.bips.coffee>
Support: Contact Lukas I guess?

This is the HEAD NODE -- for login, package installation, and job submission.

Do NOT run computations here.

- Use **salloc** for interactive sessions or
- Use **sbatch** for batch jobs

Run **slurm_user_help** to see available Slurm command aliases

If you don't know what any of this means, start with the docs linked above

Software is managed differently

Environment modules

13

- Common on HPC systems: Hundreds of users need dozens of programs
- Programs exist in dozens of versions (R, Python, specialized tools)
- Different users need different versions of different things

Solution: Environment modules

- Log in on head node: R not available
- Run `module load R/4.5.3` → R v4.5.3 is available
- Load modules **before** running `salloc` or `sbatch` — they inherit your environment
- (Python: Use `uv`, no need for modules: (docs.astral.sh/uv))

How are resources allocated?

Slurm manages resources on the cluster

- Example: You need **20 threads** for **6 hours** and **4GB of RAM** for analysis.R

Interactive: salloc

```
salloc --cpus-per-task=20 --mem=4G --time=06:00:00
```

```
salloc: Granted job allocation 137031
```

```
salloc: Nodes node01 are ready for job
```

```
# start R session
```

```
R
```

```
# Do work or run script
```

```
ranger::ranger(foo ~ ., data = bar)
```

```
source("analysis.R")
```

Batch job: sbatch

- Write an analysis.sh wrapper script

```
#!/bin/bash
```

```
#SBATCH --job-name=my-analysis
```

```
#SBATCH --output=logs/%x_%j.out
```

```
#SBATCH --cpus-per-task=20
```

```
#SBATCH --mem=4G
```

```
#SBATCH --time=06:00:00
```

```
Rscript analysis.R
```

- Run sbatch analysis.sh and go home

Slurm resource allocation

Partitions and QoS

- Slurm groups hardware in *partitions*
- 2 partitions: `compute` (12 nodes, *default*) and `gpu` (1 node)
- Prioritization is done in *Quality of Service* (QoS) queues:
- Different queues have different constraints (time, how many jobs per user)
 - **interactive**: 3 days, 2 per user (auto-applies to `salloc`)
- For `sbatch` or `#SBATCH` `--qos=<name>`:

short: 1 hour

medium (default): 1 day

long: 7 days

Slurm resource allocation

A job is not a job



16

- Slurm is very generic and supports diverse types of jobs
- Slurm also has concept of “tasks”: For us **1 task == 1 job**
- Easiest use: 1 job = 1 R/py process = Some number of threads on 1 node
- Technically possible to have 1 job spanning multiple nodes → headache country
- Slurm allocates **cores**, not **threads**
- But: You only care about threads
- You need 1 thread, you ask for `cpus-per-task=1`, you get **1** core reserved
- → More efficient to just ask for 2 “cpus” and 1 full core

Example: Efficiently using a node

If your goal is to fully utilize 1 node (96 cores, 192 threads)

What works

- 1 job with 192 threads ✓
- 96 jobs with 2 threads ✓
- 8 jobs with 24 threads ✓

What does not work

- 192 jobs with 1 thread ✗
- Each job reserves 1 **core**
- → 192 jobs allocate 192 cores, but a node only has 96!

⚠ If you're used to `batchtools`, you probably default to single-threaded jobs

Did it work?

Checking on your jobs

While running:

- `squeue --me`: list your jobs
- `tail -f logs/job_12345.out`: watch output
- `scancel <jobid>`: cancel a job

After completion:

- Check your output files (e.g. `logs/`)
- `sacct --starttime=today`: job history
- `sacct -j <jobid> --format=JobID,State,Elapsed,MaxRSS`

Common failure modes

- **TIMEOUT**: job ran out of time, request more with `--time`
- **OUT_OF_MEMORY**: request more RAM with `--mem`
- **FAILED**: your script has a bug, check the log output (*did you load modules?*)

Where does your stuff live?

Two storage tiers

Home directory

/srv/home/<user>

- Fast
- Limited space, also shared with software
- Scripts, code, *active* projects
- Both are **shared across all nodes** (head + compute) via internal network
- Compute nodes also have fast **local scratch** (\$TMPDIR, auto-cleaned after job)

Archive storage

/mnt/sas/users/<user>

- Large capacity
- Slower
- Big datasets, *inactive* projects

⚠ There are **no backups**. Use `git` for code, be careful with data.

User utilities

Command-line aliases/functions



20

- I have prepared a few shorthand tools based on things I need often
- On head node, get info with `slurm_user_help`:

```
> slurm_user_help
=== Slurm User Helper Functions ===
```

JOB MONITORING:

```
sq          - Enhanced squeue with better formatting
sqm         - My jobs queue
sqs         - Job status summary with colors
[...]
```

JOB ACCOUNTING:

```
slac        - Enhanced sacct
slac1h/1d/3d/1w - Jobs from last hour/day/3days/week (excludes PENDING)
  State filters: -a (all), -p (pending), -c (completed), -f (failed), -r (running)
sltoday     - Jobs submitted today
```

Getting started

And further reading material



21

1. Ask me for an account, I will send you login credentials
2. Bookmark the dashboard (current cluster usage):
 - <http://srvcluster.bips.de/> (no https!)
3. Read (and bookmark) the docs:
 - <http://srvcluster.bips.de/docs>
4. Try demos / usage examples for batchtools, mirai, targets:
 - srvgit.bips.eu/bips/bips-cluster-demos